Redis

* Inmemory: Caching, Rate limiting
* Traditional: Last eg DB (Inmemory: Fastest access), Remaining: Disk
* Redis: Completely on RAM (Not on HDD/SSD)
* Key-Valued DB, powerful.
* eg: Stateless app (Central Storage system), Share data by multiple apps (bruteforce attack Info): Block them at every Service/app!
* Redis: Connect disconnected/loosely conn. Services share a Common DB.

* Key-value
* Column
* Document
* Graph

Redis cli:

* SET orgID '625721321'
* GET orgID
* Don't overwrite: SETNX orgID '625721322' ⟶ No effect as already there!
* multiple sets:

```
MSET Key1 'value1' Key2 'value2' Key3 'value3'
MGET Key1 Key2 Key3 Key4
```

* Delete a Key:

```
DEL Key1
GET Key1 ⟶ #null
```

* INCR, DECR:

```
SET number 10
INCR number ⟶ #11
DECR number ⟶ #10
```

* Expire:

```
SET number 1
EXPIRE number 1 ⟶ after 1 second number is deleted
```

* TTL: Time to live:

```
SET number 1
TTL number (-1 : Key exists, no expiration set)
           (-2 : Key not exists)
           Returns expiry time remaining in sec.
```

SET and EXPIRE:

SETEX Key 10 'value'

Redis Info:

* Single threaded process : even on multicore system
* Intended to safeguard read/writes!
* If multithreaded: Need locking (Performance will be worse)
* Supports append only file persistence.

Blocks:

```
MULTI
SET hello 'world'
INCR number
EXPIRE hello 10
EXEC
```

→ After this Run everything at once.

Sharding: Distribute dataset across multiple Redis instances

Redis:

* OpenSource, BSD licensed, Advanced Key-Value Store
* Data Structure Server (String, hashes, list, Sets..)
* Written in C
* High Performance, Scalable web apps → In memory (use disk only for persistence)
  - Rich data types
  - Replicate to any no. of slaves.
* Very fast
  - 110K SETS
  - 81K GETS
* All Redis operations are atomic (update value will be received)
* Multi utility tool: Cache, messaging queue (AS Redis supports Subscribe/publish)
* Any short lived data (web app sessions, hit count)

Redis Config:

* CONFIG GET *
* CONFIG GET loglevel
* CONFIG SET loglevel 'SEVERE'

Data types:

1. String: (512 Mb)

SET name 'tutorialspoint'

GET name

2. Hashes:

* Collection of Key-value pairs (map b/w String fields & String values): represent objects.

* HMSET Key 'username' 'password' 'hello' 'hi' 'how'

HMSET.

* HGETALL Key

Store user objects : upto $2^{32} - 1$ field value pairs (more than 4 billion)

> HMSET hello hi how are you

> HGETALL hello

'hi'

'how'

'are'

'you'

3. List (List of Strings: sorted by insertion order)

> LPush JavaTpoint java

> LPush JavaTpoint mongo

> LRange JavaTpoint 0 10

'mongo'

'redis'

> LPUSH users 1 2 3 4

> LINDEX users 0 ⟶ #1

> LRANGE users 0 -1 (slice)

$2^{32} - 1$ elements

4. Sets

*Unordered Collection of String :- O(1) add, Remove, Search for existence.

> Sadd Key1 redis

> Sadd Key1 mongo

> Sadd Key1 mongo

> Smembers Key1 ⟶ 'mongo'
'redis'

5. Sorted sets:

* Non-repeating collection of String: Sorted. (say small → great)
  → Score

```
> zadd  key1  0  redis
> zadd  key1  0  mongo
> zadd  key1  0  rabbit
> ZRANGEBYSCORE  key1  0  1000
  "redis"
  "mongo"
  "rabbit"
```

When to use NoSql databases:

* Large amount of data:
* Highly replicable: No need for primary read/write & only secondary read-only nodes. **Consistent**

Horizontal Scaling: More machines
Vertical Scaling: More CPU, RAM

* While horizontal scaling: RDBMS takes effort!

Sharding: Split/partition resources into smaller pieces & distribute to different computing resources. [Use closest DB: ↓ latency]

* Read/write: higher in NoSql (unstructured data)
* Traditional DB: Takes effort to change schema
* Resource intensive app: don't scale well.

NOSQL: Scale at ease!

Redis:

* Data storage (Durability guaranteed by Redis)
* Data Structures → String, List, Set, Sorted Set, Hash, Bit array, Hyperlog logs, Streams, Geospatial indexes
* Supports publish/subscribe pattern.

Multimodel databases:

* Single integrated backend server app
* Packs modules: pre built (JSON support, SQL, Image processing, Linear algebra, Indexes)

Why not single model:

* Bottleneck in access & represent data!
* multimodel databases handles
  1. Relational
  2. object-oriented
  3. Key value
  4. Wide column
  5. Document
  6. Graph models.

* Store structured & semi structured data
* Eliminates fragmentation problem.

Top multimodel DB stores:

- → Data storage, backup, End recovery
- → Query, Index (query, use index: efficient query)
- → ACID (compliant, fault tolerant)
- → Integration, Advanced security.

Usecase:

* Auto complete, Result highlight
* Real time analysis: Top score, Cost, post, bidding
* Fraud detection: Spot trends
* Gaming & leaderboards
* Session management
* Social apps
* Recommendation manager

* Cache, publish/subscribe pattern for incoming data
* Job/Queue management
* Built in analysis
* Native JSON handling.

Search:

* Index, query : Search with high performances
* Secondary Indexes

Cache: B/w Server & app (Freeup dB)

* No DB, No tables
* SET ⟶ To create data
* Writes to disk at varying time interval (durability incase of failures)

---

* redis-cli -h <hostname> -P <port>

* Keys can be anything (As Binary safe) : Also use an image as key.
* But mostly strings (common)

Hashes:

HSET house:5100 numBed 3 size 6000 hvac 'forced'

HGET house:5100 numBed ⟶ #3

Sorted Sets:

* Leaderboard (score : each member)

> ZADD users 31 steve 2 owen 13 jake
> ZRANGE user:followers 0 -1 (0 to end.)
> ZRANGE users 0 -1 WITHSCORES
> ZREVRANGE users 0 -1 WITHSCORES
> ZINCRBY users 20 jakob ⟶ Increment jakob by 20 # 33

Hyperlog:

* Keep an estimate count of unique items (eg: Track Count of unique visitors to a website). maintains internal hash. (determine already there : If yes : Not entered.)

> PFADD visitors 127.0.0.1 ⟶ 1 (new)
0 (If already exists)

> PFCOUNT ⟶ No. of unique hyperlogs.

Pub/Sub:

* Redis can act as a fast & efficient means to exchange messages in a publisher/subscriber pattern.

* publisher creates key value pair : 0/more subscribers to receive messages.

> PUBLISH weather temp:85f

'The message is published on the channel weather'
The client subscribed will receive
'message'
'weather'
'temp:85f'

PUBLISH weather: 54481 temp: 85$

PSUBSCRIBE weather: [*]

Geospatial Indexes:

* Latitude & Longitude data (deskance)

> GEOADD towers −89.500 44.500 tower1
> GEODIST towers tower1 tower2 # calculate deskance
> GEODIST towers tower1 tower2 mi

Redis Streams:

* data is appended like a log file (So only Stream)

XADD
XRANGE

eCan view pending messages & do powerful operations'

Redis modules:

1. Redis Search: Full text search engine (with secondary Indexing)
   powerful Querying
   weighted search
2. RedisJSON: Store JSon, Inmemory manipulation.
   (product catalogs, 3-rd party feeds).
3. Redis TimeSeries: Store time Series data (Added timestamp)
4. Redis Graph: Graph DB
5. Redis Bloom: Support additional probabilistic DS
6. Redis AI
7. Redis Greans: Batch/event driven processing.

---

* Shard: Takes care of a subset of data
* proxy (zero latency): Proxy to appropriate shard (each node of Cluster uses proxy.
* Cluster manager: manages cluster health, monitoring (balance, Shard, Provision/deprovision)

A: write full / don't (No partials): Atomic
C: Data Correct (Before/after write): Consistent,
I: Each process: Seperate (Isolation)
D: Durability: Ensure data persistence:
Transaction Complete, retrieved
Incase of failure.

Redis:

A: MULTI, WATCH, EXEC [Indivisible & Irreducible]

C: only permitted writes

I: Single threaded (only multi/exec: executed)

D: Durability

Data persistence: (2 methods)

1. AOF (Append only file):

* Redis replies for each successful operation write
* Aof: applies to every shard
* write every second: fast but not safe (slower performance)

* Redis enterprise handles this differently! (writing): optimized
* performance unaffected as master shard unaffected.

2. SnapShot:

* point in time copy (For durability rather than as backup)

CAP theorem:

* Impossible for a network based service (Server) to fully provide more that 2 out of

* Consistency (Conflict free replications)
* Availability (make Copies across DCs)
* partition tolerance.

Layers:

management layer: Administer cluster, placement of Shards, Failure detection & mitigation

Data access layer: manage connections (with clients, pri/Sec Shard)

In memory replication using WAN: Sync!

Conflict free replicated Datatypes

* CRDTS: multiple Copies stored across locations (Independently)
* update, Resolve Inconsistencies.

Conflict: which one to use (math rules)

Faster, fault tolerant.

Resources:

* Time
* Bit: No. of operations on bits required to run an algorithm
* Space

Atomicity: Either occurs/not! (MULTI, EXEC) → errors → Unable to queue (before starting) / Error in queuing!

* EXEC: All done, All failed.
* Redis: doesn't support rollbacks!

WATCH mykey:

* Conditional 'EXEC': perform only when watched keys are unmodified.
* modification: Client, Redis (expiration/eviction)
* when EXEC is called: all keys are unwatched!

'Optimized locking'

Inmemory DB:

* Relies on RAM (Fewer CPU instruction), Volatile.
* Split data into multiple Redis instances.

Redis persistence:

* Redis DB: Snapshot in deterministic time interval
* AOF: write in logs (every operation): Do everything to restore!

RDB → perfect for backup (everyday to 30days)
→ disaster recovery (DC), Faster restart!
→ Not good to backup time long (data lost high as interval ↑)
→ Time consuming if done often affects performance.

AOF → Background (done), Append to old file
→ Bigger files than RDB! As fast as RDB
→ memory usage high!

Snapshot

Save 60 1000 [every 60s, 1000 keys changed atleast]